

# CS-541: Artificial Intelligence

## Lecture 5

Abdul Rafae Khan

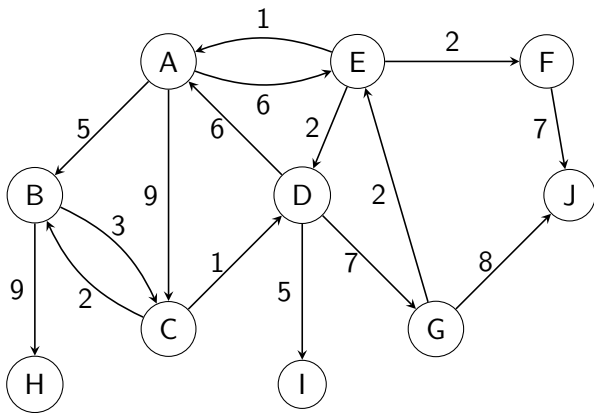
Department of Computer Science  
Stevens Institute of Technology  
*akhan4@stevens.edu*

February 22, 2022

# Uniform Cost Search

start state: A

goal state: H, I, J



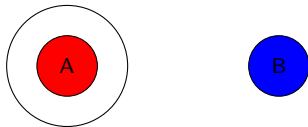
# Uniform Cost Search

---



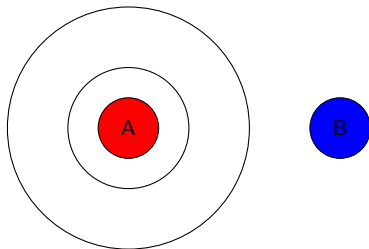
# Uniform Cost Search

---



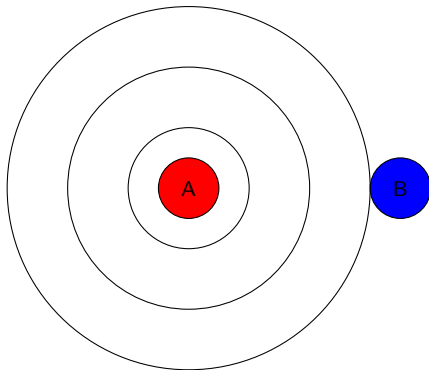
# Uniform Cost Search

---



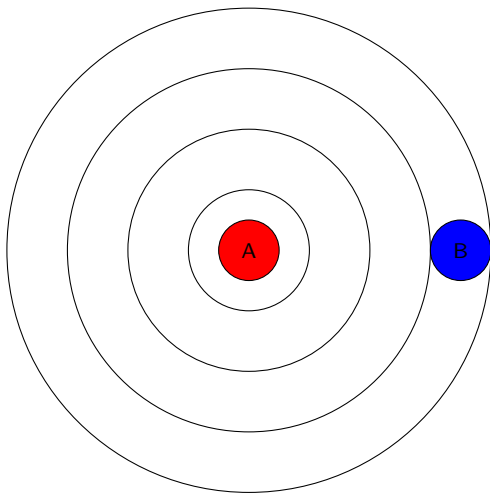
# Uniform Cost Search

---



# Uniform Cost Search

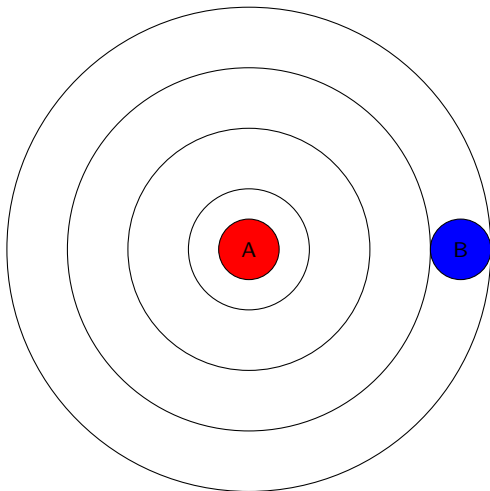
---



# Uniform Cost Search

---

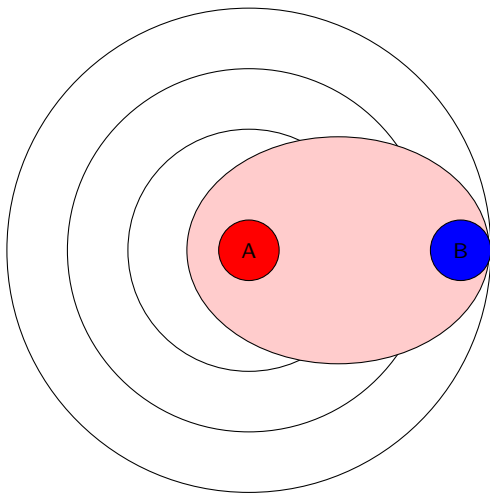
A lot of Wasted effort!





# Informed Search

---



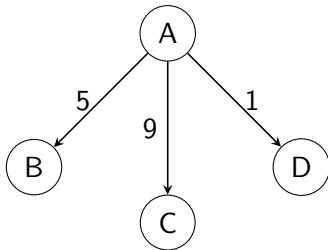
# Informed Search

---

## Uniform cost search:

Search using:  $\min_{a \in \text{Actions}(s)} \text{Cost}(s, a)$

At every state  $s$ , select the action  $a$  with the minimum cost



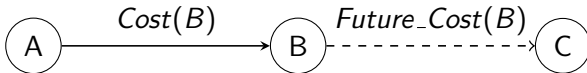
# Informed Search

---

$$Total\_Cost(goal) = Cost(s, a) + Future\_Cost(s)$$

$Future\_Cost(s)$  is the cost to go from state  $s$  to state  $goal$

If we know the  $Future\_Cost(s)$ , then our search would be better



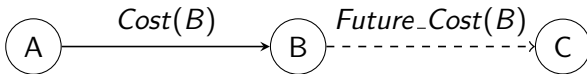
# Informed Search

---

## Informed cost search:

Search using:  $\min_{a \in \text{Actions}(s)} \text{Cost}(s, a) + \text{Future\_Cost}(s)$

At every state  $s$ , select the action  $a$  with the minimum cost



# Informed Search

---

If we know  $Future\_Cost(s)$ , we know the solution!

# Informed Search

---

Lets estimate the *Future\_Cost*(*s*)

$$Cost'(s) = Cost(s) + h(Successor(s)) - h(s)$$

*h*(*s*) is an estimate of the cost from *s* to the *goal*

# A\* Search

---

**Start:** D

**Goal:** G



*Whiteboard*

# A\* Search

---

**Start:** D

**Goal:** G



Node	A	B	C	D	E	F	G
$h(s)$	6	5	4	3	2	1	0



# A\* Search

Start: D

Goal: G



Node	A	B	C	D	E	F	G
$h(s)$	6	5	4	3	2	1	0

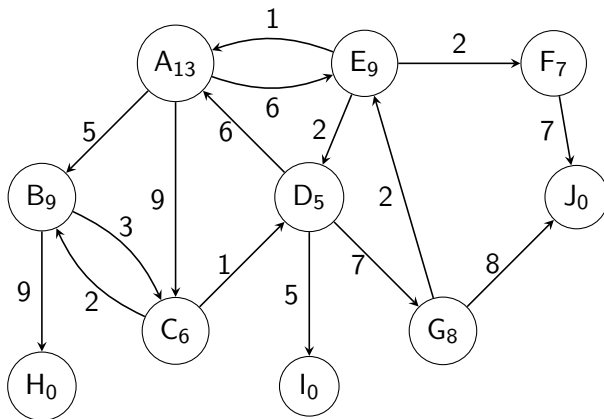
$$Cost'(C) = Cost(C) + h(C) - h(D) = 1 + 4 - 3 = 2$$

*Whiteboard*

# A\* Search

start state: A

goal state: H, I, J



Whiteboard

# Tree Search

---

Can any heuristic work?

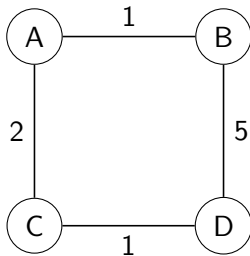
# Tree Search

---

Can any heuristic work?

**start state:** A

**goal state:** D



*Whiteboard*

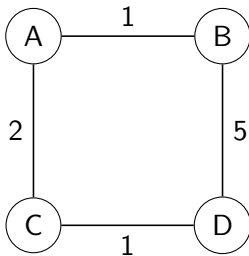
# Tree Search

---

Can any heuristic work?

**start state:** A

**goal state:** D



Node	A	B	C	D
$h(s)$	0	0	100	0

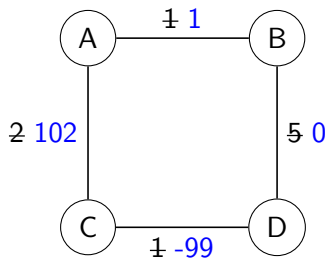
*Whiteboard*

# Tree Search

Can any heuristic work?

**start state:** A

**goal state:** D



Node	A	B	C	D
$h(s)$	1	1	100	0

*Whiteboard*

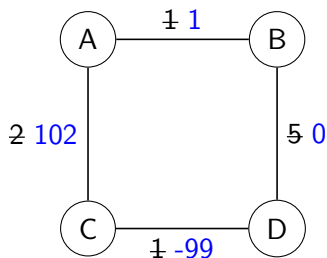
# Tree Search

---

Can any heuristic work?

**start state:** A

**goal state:** D



The algorithm does not work with **negative edge costs**

*Whiteboard*

# Conditions for heuristics

---

## Consistency

1.  $h(\text{successor}(s)) - h(s) \leq \text{Cost}(s, \text{successor}(s))$

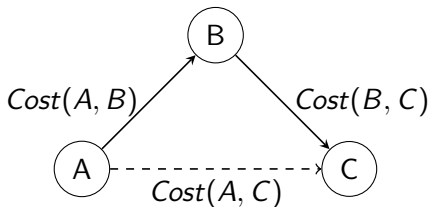


# Conditions for heuristics

---

## Consistency

1.  $h(\text{successor}(s)) - h(s) \leq \text{Cost}(s, \text{successor}(s))$



$$h(A) - h(C) \leq \text{Cost}(A, C)$$

or  $h(A) \leq \text{Cost}(A, C) + h(C)$

# Conditions for heuristics

---

## Consistency

1.  $h(\text{successor}(s)) - h(s) \leq \text{Cost}(s, \text{successor}(s))$
2.  $h(\text{goal}) = 0$

# Conditions for heuristics

---

## Consistency

1.  $h(\text{successor}(s)) - h(s) \leq \text{Cost}(s, \text{successor}(s))$
2.  $h(\text{goal}) = 0$

## Admissibility

1.  $h(s) \leq \text{Cost}(s, \text{goal})$

# Conditions for heuristics

---

## Consistency

1.  $h(\text{successor}(s)) - h(s) \leq \text{Cost}(s, \text{successor}(s))$
2.  $\text{Future\_Cost}(\text{goal}) = 0$

## Admissibility

1.  $h(s) \leq \text{Cost}(s, \text{goal})$ , heuristic cost is less than actual cost
- $$h(A) \leq \text{Cost}(A, C) + h(C) \implies g(A) + h(A) \leq g(A) + \text{Cost}(A, C) + h(C)$$

# Conditions for heuristics

---

## Consistency

1.  $h(\text{successor}(s)) - h(s) \leq \text{Cost}(s, \text{successor}(s))$
2.  $\text{Future\_Cost}(\text{goal}) = 0$

## Admissibility

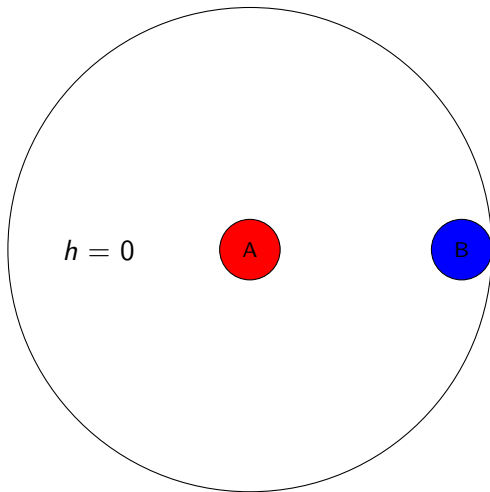
1.  $h(s) \leq \text{Cost}(s, \text{goal})$ , heuristic cost is less than actual cost
- $$h(A) \leq \text{Cost}(A, C) + h(C) \implies g(A) + h(A) \leq g(A) + \text{Cost}(A, C) + h(C)$$

$h(s)$  is consistent  $\implies h(s)$  is admissible

# Informed Search

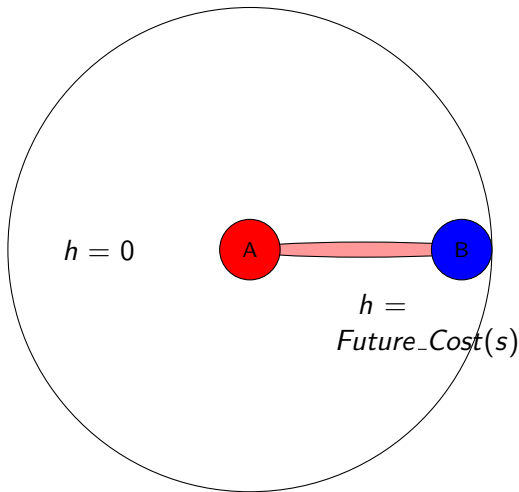
---

If  $h(s) = 0$ ,  $A^*$  is the same as UCS



# Informed Search

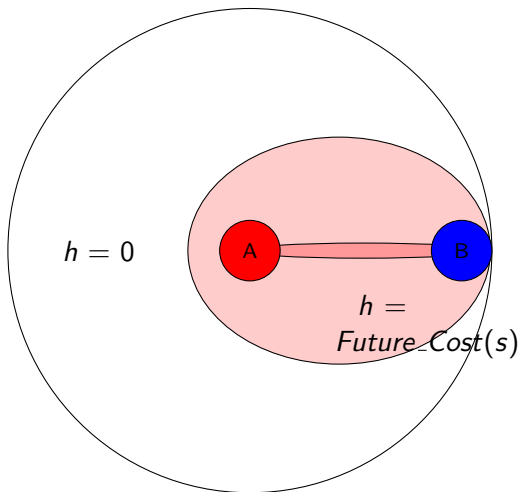
If  $h(s) = \text{Future\_Cost}(s)$ , the  $A^*$  only explores the minimum cost nodes



# Informed Search

---

Usually  $h(s)$  is in between





# Heuristics

---

How to get heuristic values?

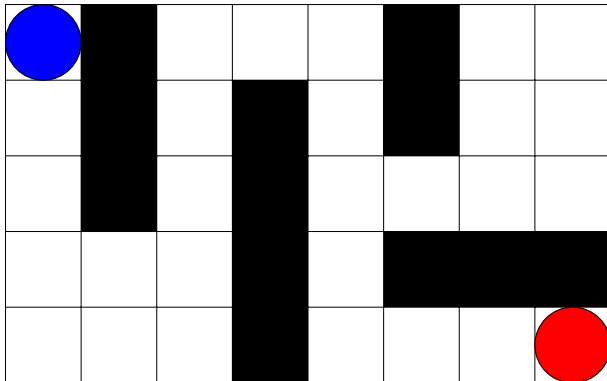
# Relaxations

---

How to get heuristic values? Make the problem easy!

# Relaxations

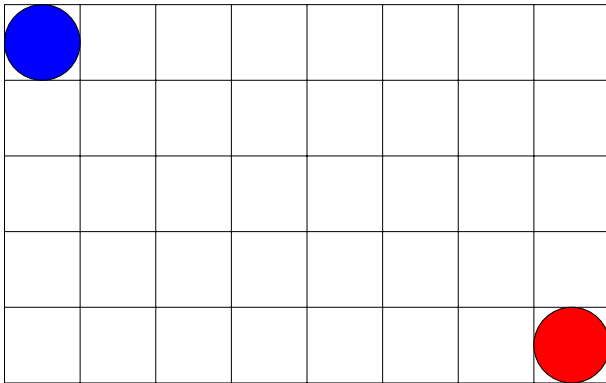
---



# Relaxations

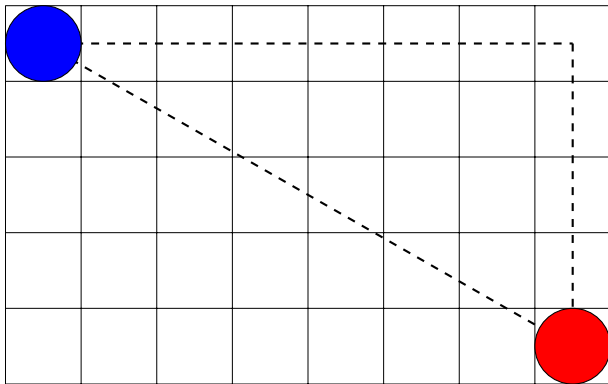
---

Drop all the edges!



# Relaxations

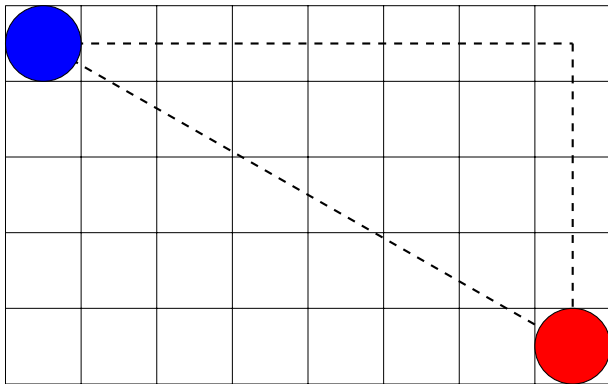
---



- $h(\text{start}) = \text{Manhattan}(\text{start}, \text{goal})$
- $h(\text{start}) = \text{Euclidean}(\text{start}, \text{goal})$

# Relaxations

---



- $Cost(start, right) = \infty$
- $Cost'(start, right) = \mathbb{R}^+$

# Relaxations

---



Goat & Cabbage/Goat & Wolf can be left alone

# Relaxations

---



Goat & Cabbage/Goat & Wolf can be left alone  
Similarly for other problems!



# Local Search

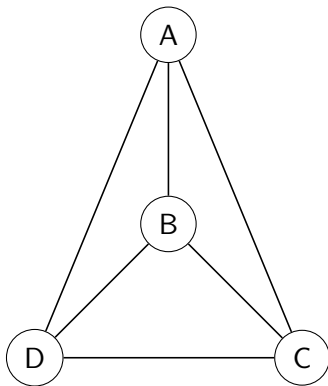
---

Sometimes the path is not important  
The goal is the solution we want to achieve!

# Traveling Salesman Problem

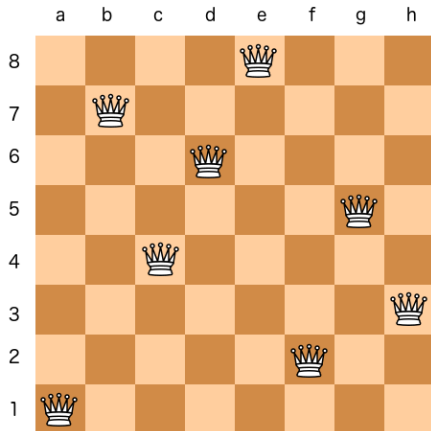
---

Start at state *A* and reach state *C*  
while visiting all the states exactly once



# 8-Queen Problem

---



Place 8 queens on a chessboard such that no queen is being attacked

# Local Search

---

Sometimes the path is not important  
The goal is the solution we want to achieve!

Do not require a lot of memory  
Work with continuous space

# Local Search

---

Sometimes the path is not important  
The goal is the solution we want to achieve!

Do not require a lot of memory  
Work with continuous space

You can think of it as optimizing a criterion  
You can iteratively try to improve the current state

# Hill Climbing

---

Start at any state

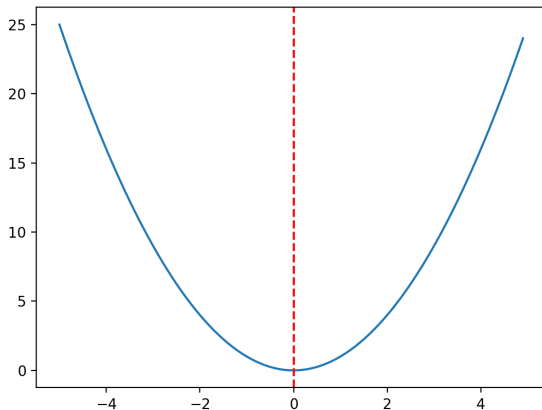
Repeatedly move to the best neighboring state

If no neighboring state is better, stop

# Local vs Global Maxima/Minima

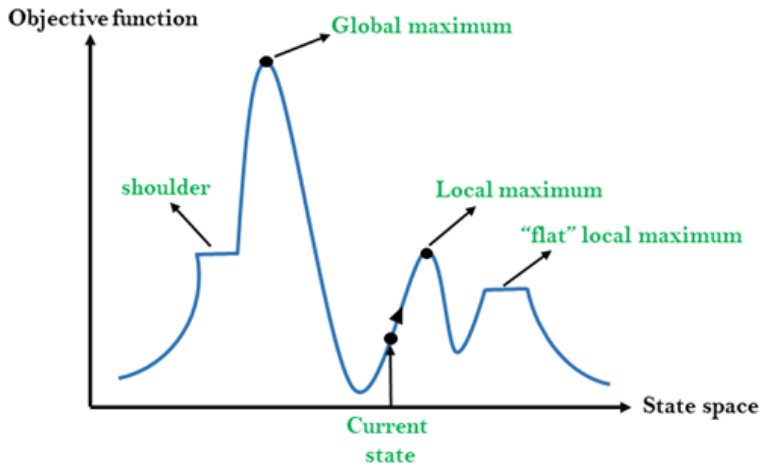
---

Maximize the benefit or minimize the cost



# Local vs Global Maxima/Minima

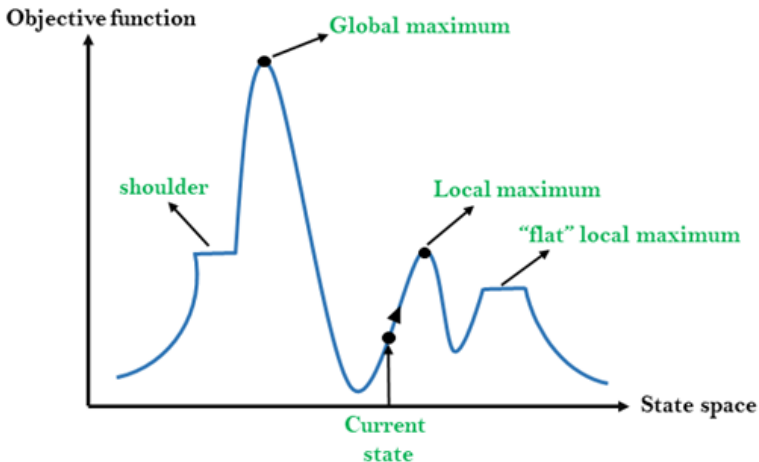
---





# Local vs Global Maxima/Minima

Restart hill climbing from a random state



# Beam Search

---

Lets say we have a machine translation model already trained

Translation output is one word at a time

We want to use it to output the best translation for any given input sentence



# Beam Search

---

Given the input sentence (french) and the previous output word (English), the model outputs the next word



# Beam Search

---

If we have  $V$  words in the vocabulary and maximum sentence length is  $L$

Exhaustive search will be  $V^L$

If  $vocab = 1000$  and  $max\_length = 10$ , then  $10^{40}$  possible choices to select from



# Beam Search

---

A simpler option is to select the best at each position



# Beam Search

---

Typically greedy selection for the word at each position

Position	1	2	3	4
A	0.5	0.1	0.2	0.1
B	0.2	<b>0.4</b>	0.2	0.2
C	0.2	0.3	0.4	0.1
<END>	0.1	0.2	0.1	0.6

Output probability =  $0.5 \times 0.4 \times 0.4 \times 0.6 = 0.048$

# Beam Search

---

It can happen that this is not the optimal

Selecting *2nd* best at position 2 gives better total probability

Position	1	2	3	4
A	0.5	0.1	0.1	0.1
B	0.2	0.4	0.6	0.2
C	0.2	<b>0.3</b>	0.2	0.1
<END>	0.1	0.2	0.1	0.6

Output probability =  $0.5 \times 0.3 \times 0.6 \times 0.6 = 0.054$

# Beam Search

---

1st step:

Select the  $K$  maximum probability words

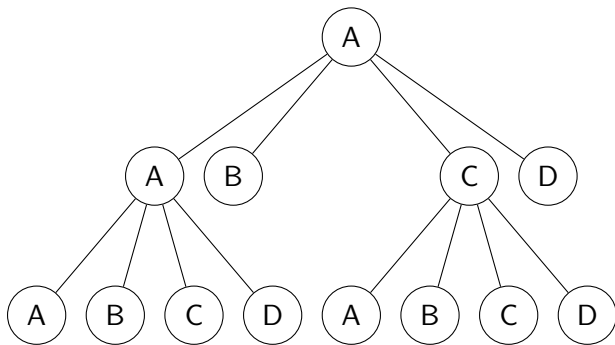
For each subsequent step:

Fix the previous selections and generate  $K$  possibilities and select the overall  $K$  maximum



# Beam Search

---



*Whiteboard*

# Games

---

All previous algorithms work for single agent only  
What if we have more than one agent?

# Games

---

All previous algorithms work for single agent only  
What if we have more than one agent?

- Tic-Tac-Toe
- Pacman
- Chess

# Games

---

Assumptions:

- States and transitions are deterministic
- Environment is fully observable
- 2 agents with turn-taking
- Zero-sum game: In which result is an advantage for one side and a loss for the other (One winner & One loser)

Algorithm to compute the strategy for to recommend every possible move for

# Tree Search

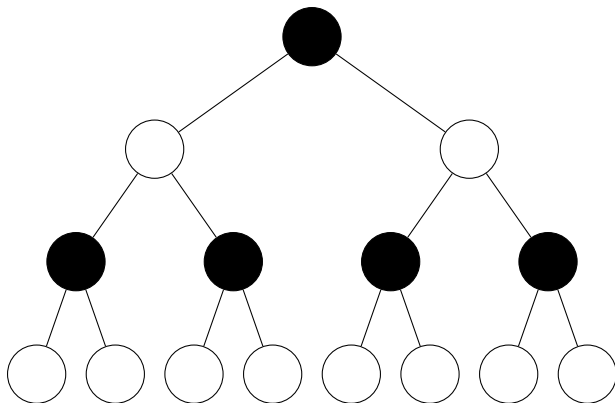
---

Each node represents a turn by a player

One player tries to maximize a utility and the other tries to minimize

# Minmax Search

---



# Search Tree

---

Each node represents a turn by a player

One player tries to maximize a utility and the other tries to minimize

Leaf nodes represent some utility value

Root to leaf is a path for the outcome of the game

Each player's turn is called a ply

One turn is the play by both the players

# Search Tree

---

Leaf nodes can represent a utility value

What can be possible values for

Tic-Tac-Toe

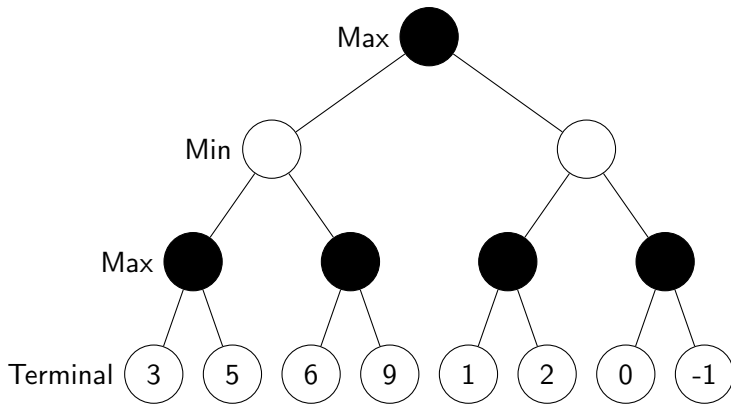
Chess

Pacman



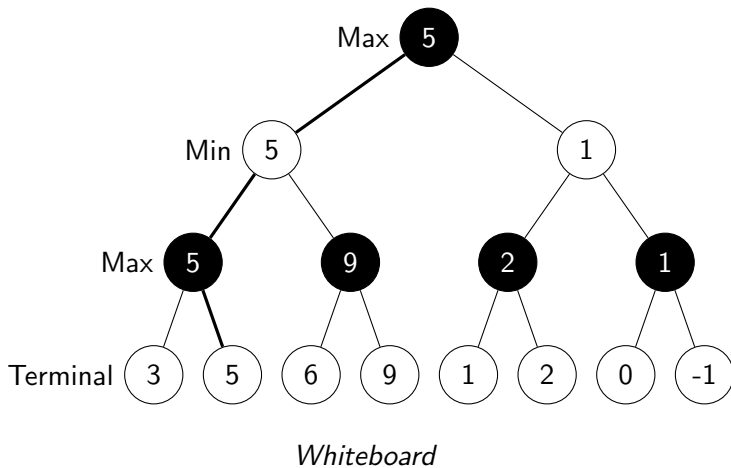
# Search Tree

---



# Minmax Search

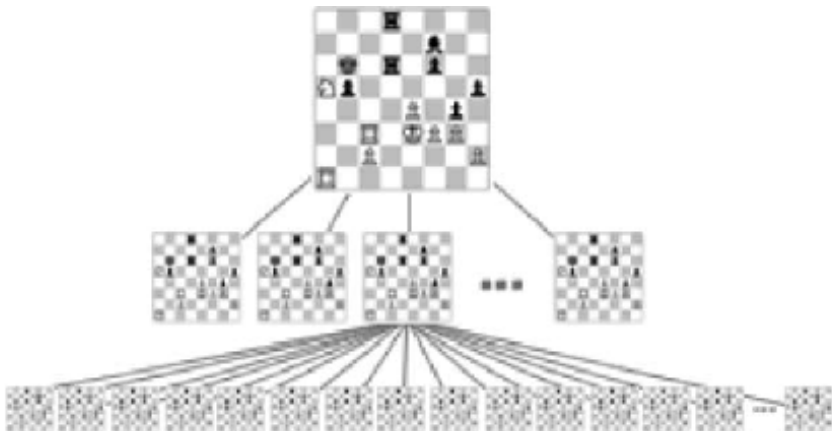
---



# Alpha-Beta Pruning

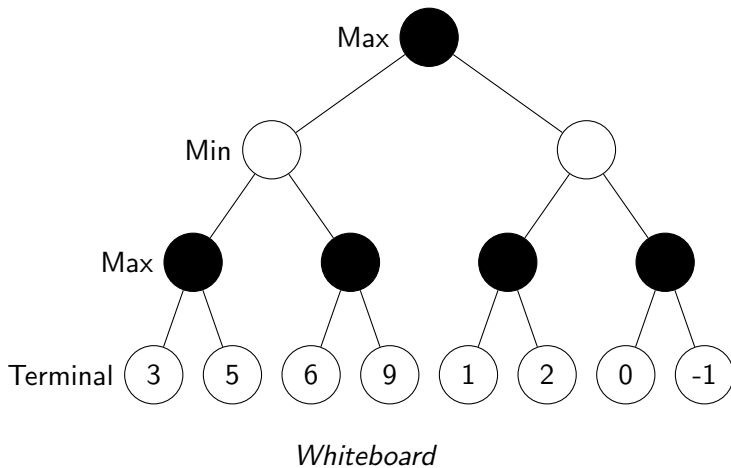
The complete tree can be very big

E.g. Just the first 5 turns (10 plies) of chess generate **69,352,859,712,417** possibilities



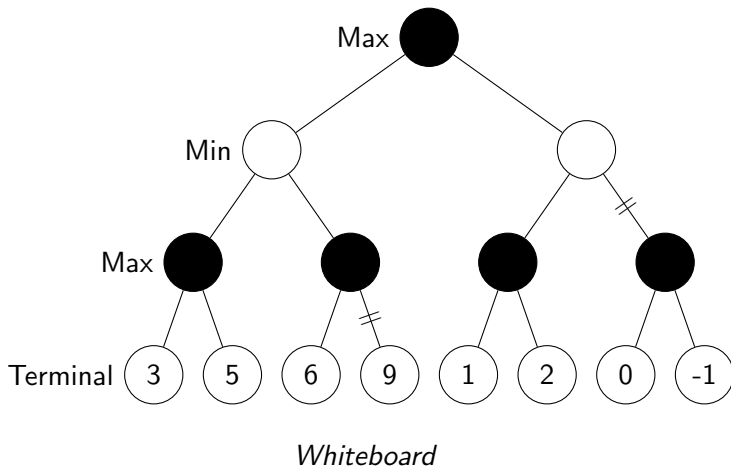
# Alpha-Beta Pruning

---



# Alpha-Beta Pruning

---



# Alpha Beta Pruning

---

```
def max_value(s,  $\alpha$ ,  $\beta$ ):  
    if leaf_node(s):  
        return value(s)  
    v =  $-\infty$   
    for each u in successor(s):  
        v' = max(v, min_value(u,  $\alpha$ ,  $\beta$ ))  
        if v'  $\geq$   $\beta$ :  
            return v  
         $\alpha$  = min(v,  $\alpha$ )  
    return v
```

```
def min_value(s,  $\alpha$ ,  $\beta$ ):  
    if leaf_node(s):  
        return value(s)  
    v =  $+\infty$   
    for each u in successor(s):  
        v' = min(v, max_value(u,  $\alpha$ ,  $\beta$ ))  
        if v'  $\leq$   $\alpha$ :  
            return v  
         $\beta$  = min(v,  $\beta$ )  
    return v
```

# Recap

---

## Informed Search

- A\* Search
- Heuristics

## Local Search

- Hill Climbing
- Beam Search

## Zero-Sum Games

- Minimax Search
- Alpha-Beta Pruning

# References

---



Stuart Russell and Xiaodong Song (2021)

CS 188 — Introduction to Artificial Intelligence

*University of California, Berkeley*



Chelsea Finn and Nima Anari (2021)

CS221 — Artificial Intelligence: Principles and Techniques

*Stanford University*



# The End