CS 559: Machine Learning Fundamentals and Applications

Lecture 7

Lecturer: Xinchao Wang

xinchao.wang@stevens.edu

Teaching Assistant: Yiding Yang

yyang99@stevens.edu

Outline

- Linear Regression
  - Regression technique
- Logistic Regression (called regression, but in fact, classification)
  - Classification technique
- Gradient Descent/Ascent
  - Descent -> min.
  - Ascent -> max.
  - Optimization technique
- Perceptron
  - Classification technique

Linear Regression

Simple Linear Regression

- How does a single variable of interest relate to another (single) variable?
  - Y = outcome variable (response, dependent...)
  - X = explanatory variable (predictor, feature, independent...)
- Data: n pairs of continuous observations (X_1,Y_1) … (X_n,Y_n)

Example

- How does systolic blood pressure (SBP) relate to age?

- Graph suggests that Y relates to X in an approximately linear way

Regression: Step by Step

1. Assume a linear model: $Y = \beta_0 + \beta_1 X$
2. Find the line which “best” fits the data, i.e. estimate parameters β_0 and β_1
3. Does variation in X help describe variation in Y ?
4. Check assumptions of model
5. Draw inferences and make predictions

Straight-line Plots

Assumptions of Linear Regression

- Five basic assumptions

1. Existence: for each fixed value of X, Y is a random variable with finite mean and variance
2. Independence: the set of Y_i are independent random variables given X_i

Assumptions of Linear Regression

3. Linearity: the mean value of Y is a linear function of X

$$\mu_{Y|X} = E[Y|X] = \beta_0 + \beta_1 X$$

Assumptions of Linear Regression

4. Homoscedasticity: the variance of Y is the same for any X
5. Normality: For each fixed value of X, Y has a normal distribution (by assumption 4, σ^2 does not depend on X)

$$Y \sim N(\mu_{Y|X}, \sigma^2)$$

Formulation

$$Y_i = \beta_0 + \beta_1 X_i + \varepsilon_i$$

$$\varepsilon_i \sim N(0,\sigma^2) \quad \text{independent}$$

$$\Rightarrow Y_i \text{ are } N(\beta_0 + \beta_1 X_i, \sigma^2) \text{ given } X_i$$

$$\Rightarrow E(Y_i \mid X_i) = \beta_0 + \beta_1 X_i$$

$$\text{Var}(Y_i \mid X_i) = \sigma^2 = \text{variability of } Y_i \text{ about } \mu_{Y \mid X_i}$$

- Y_i are linear function of X_i plus some random error

$$\varepsilon_i = \text{error } \left(= Y_i - \mu_{Y \mid X_i}\right)$$

$$\hat{\varepsilon}_i = Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i \quad \text{is called the "residual" for } Y_i$$

$$\hat{\varepsilon}_i = Y_i - \hat{Y}_i$$

Linear Regression

Estimating β_0 and β_1

- Find "best" line
- Criterion for "best": estimate β_0 and β_1 to minimize:

$$\sum_{i=1}^{n}\left(Y_i - \hat{\beta}_0 - \hat{\beta}_1 X_i\right)^2$$
$$= \sum_{i=1}^{n}\left(Y_i - \hat{Y}_i\right)^2 = \sum_{i=1}^{n} \hat{\varepsilon}_i^2$$

- This is the residual sum of squares, sum of squares due to error, or sum of squares about regression line
- Least Squares estimator

Rationale for LS Estimates

- $\hat{\varepsilon}^2$ measures the "deviance" of Y_i from the estimated model
- The "best" model is the one from which the data deviate the least

Least Squares Estimators

- Taking derivatives with respect to β, we obtain

$$\hat{\beta}_1 = \frac{\sum_{i=1}^{n}\left(X_i - \bar{X}\right)\left(Y_i - \bar{Y}\right)}{\sum_{i=1}^{n}\left(X_i - \bar{X}\right)^2} = \frac{S_{xy}}{S_x^2}$$

$$\hat{\beta}_0 = \bar{Y} - \hat{\beta}_1 \bar{X}$$

Example: SBP/age data

$$\bar{X} = 45.13 \quad \bar{Y} = 142.53$$

$$S_x^2 = 233.91 \quad S_y^2 = 509.91 \quad S_{xy} = 227.10$$

$$\hat{\beta}_1 = \frac{s_{xy}}{s_x^2} = \frac{227.10}{233.91} = 0.97$$

$$\hat{\beta}_0 = \bar{Y} - \hat{\beta}_1 \bar{x} = 142.53 - 0.97(45.13) = 98.71$$

$$\hat{y} = 98.71 + .97x$$

.97mm Hg ↑ for every 1 yr ↑ in age

Using the Model

- Using the parameter estimates, our best guess for any Y given X is

$$\hat{Y} = \hat{\beta}_0 + \hat{\beta}_1 X$$

- Hence at $\bar{X}$

$$\hat{\beta}_0 + \hat{\beta}_1 \bar{X} = \left(\bar{Y} - \hat{\beta}_1 \bar{X}\right) + \hat{\beta}_1 \bar{X} = \bar{Y}$$

- Also $\hat{Y} = \bar{Y} + \hat{\beta}_1 \left(X - \bar{X}\right)$

Regression and Correlation Coefficient

Regression Coefficient

Correlation Coefficient

$$\hat{\beta}_1 = \frac{S_{xy}}{S_x^2} \text{ and } r_{xy} = \frac{S_{xy}}{S_x S_y}$$

$$\Rightarrow r_{xy} = \hat{\beta}_1 \cdot \frac{S_x}{S_y}$$

Example

Suppose we have:

x	y
1	2
2	4
-1	0
3	4
4	6

$$n = 5$$

$$\bar{x} = \frac{1+2-1+3+4}{5} = 1.8$$

$$\bar{y} = \frac{2+4+0+4+6}{5} = 3.2$$

Calculate r:

$$r = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum (x_i - \bar{x})^2 \sum (y_i - \bar{y})^2}}$$

$x_i - \bar{x}$	$y_i - \bar{y}$	$(x_i - \bar{x})(y_i - \bar{y})$
1-1.8 = -0.8	2-3.2 = -1.2	(-0.8)(-1.2) = 0.96
2-1.8 = 0.2	4-3.2 = 0.8	(0.2)(0.8) = 0.16
-1-1.8 = -2.8	0-3.2 = -3.2	(-2.8)(-3.2) = 8.96
3-1.8 = 1.2	4-3.2 = 0.8	(1.2)(0.8) = 0.96
4-1.8 = 2.2	6-3.2 = 2.8	(2.2)(2.8) = 6.16
0	0	17.2

Example

$$\sum (x_i - \bar{x})^2 = (-0.8)^2 + (0.2)^2 + (-2.8)^2 + (1.2)^2 + (2.2)^2 = 14.8$$

$$\sum (y_i - \bar{y})^2 = (-1.2)^2 + (0.8)^2 + (-3.2)^2 + (0.8)^2 + (2.8)^2 = 20.8$$

$$S_x = \sqrt{\frac{14.8}{4}} = 1.92$$

$$s_y = \sqrt{\frac{20.8}{4}} = 2.28$$

$$r = \frac{17.2}{\sqrt{(14.8)(20.8)}} = 0.98$$

$$\hat{\beta}_1 = r \cdot \frac{s_y}{S_x} = (0.98)\left(\frac{2.28}{1.92}\right) = 1.16$$

$$\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x} = 3.2 - (1.16)(1.8) = 1.11$$

Estimated regression line is $\hat{y}_i = 1.11 + 1.16 x_i$

Example

Fitted values are the $\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i$

$\hat{y}_1 = 1.11 + 1.16(1) = 2.27$ $\hat{y}_4 = 1.11 + 1.16(3) = 4.59$

$\hat{y}_2 = 1.11 + 1.16(2) = 3.43$ $\hat{y}_5 = 1.11 + 1.16(4) = 5.75$

$\hat{y}_3 = 1.11 + 1.16(-1) = -0.05$

Residuals =

$y_i - \hat{y}_i = \hat{\varepsilon}_i$	$(y_i - \hat{y}_i)^2$
2-2.27 = -0.27	$(-0.27)^2 = 0.073$
4-3.43 = 0.57	$(0.57)^2 = 0.345$
0-(-0.05) = 0.05	$(0.05)^2 = 0.0025$
4-4.59 = -0.59	$(-0.59)^2 = 0.348$
6-5.75 = 0.25	$(0.25)^2 = 0.0625$
$\sum \hat{\varepsilon}_i = \sum (y_i - \hat{y}_i) = 0$	$\sum (y_i - \hat{y}_i)^2 = 0.811$

Example

$\hat{y} = 1.11 + 1.16x$

At $x = 3$:

$y_i - \hat{y}_i = 4 - 4.59 = -.59$

Logistic Regression

Logistic Regression

- Learn $P(Y|X)$ directly
- Consider learning $f: X \rightarrow Y$, where
  - X is a vector of real-valued features, $< X_1 \ldots X_n >$
  - Y is boolean
  - assume all X_i are conditionally independent given Y
  - model $P(X_i \mid Y = y_k)$ as Gaussian $N(\mu_{ik}, \sigma_i^2)$
  - model $P(Y)$ as Bernoulli (π)
    - Y is 1 with probability π

Derivation of P(Y|X)

$$P(Y=1|X) = \frac{P(Y=1)P(X|Y=1)}{P(Y=1)P(X|Y=1) + P(Y=0)P(X|Y=0)}$$

$$= \frac{1}{1 + \frac{P(Y=0)P(X|Y=0)}{P(Y=1)P(X|Y=1)}}$$

$$= \frac{1}{1 + \exp(\ln \frac{P(Y=0)P(X|Y=0)}{P(Y=1)P(X|Y=1)})}$$

$$= \frac{1}{1 + \exp(\ (\ln \frac{1-\pi}{\pi}) + \boxed{\sum_i \ln \frac{P(X_i|Y=0)}{P(X_i|Y=1)}})}$$

$$P(x \mid y_k) = \frac{1}{\sigma_{ik}\sqrt{2\pi}} \quad e^{\frac{-(x-\mu_{ik})^2}{2\sigma_{ik}^2}}$$

$$\boxed{\sum_i \left(\frac{\mu_{i0} - \mu_{i1}}{\sigma_i^2} X_i + \frac{\mu_{i1}^2 - \mu_{i0}^2)}{2\sigma_i^2} \right)}$$

$$P(Y=1|X) = \frac{1}{1 + \exp(w_0 + \sum_{i=1}^{n} w_i X_i)}$$

Very Convenient

$$P(Y=1|X=<X_1,...X_n>) = \frac{1}{1+exp(w_0+\sum_i w_i X_i)}$$

implies

$$P(Y=0|X=<X_1,...X_n>) = \frac{exp(w_0+\sum_i w_i X_i)}{1+exp(w_0+\sum_i w_i X_i)}$$

implies

$$\frac{P(Y=0|X)}{P(Y=1|X)} = exp(w_0+\sum_i w_i X_i)$$

implies

$$\ln\frac{P(Y=0|X)}{P(Y=1|X)} = w_0+\sum_i w_i X_i$$

linear classification rule!

Very Convenient

- Posteriors sum to 1 and remain in [0, 1]
- Logit: $l = logit(p) = \log\left(\frac{p}{1-p}\right)$=α+βx
  - l is linear in x
- Probability: $p = \frac{e^l}{1+e^l}$

Logistic Function

$$P(Y=1|X) = \frac{1}{1+\exp(w_0 + \sum_{i=1}^{n} w_i X_i)}$$

- $p = 0, \quad l = \log\left(\frac{p}{1-p}\right) = -\infty$
- $p = \frac{1}{2}, \quad l = \log\left(\frac{p}{1-p}\right) = 0$
- $p = 1, \quad l = \log\left(\frac{p}{1-p}\right) = +\infty$

Logistic Regression More Generally

- Logistic regression when Y is not boolean (but still discrete)
  - $y \in \{y_1 \ldots y_R\}$: learn R-1 sets of weights
  - for k<R $$P(Y = y_k|X) = \frac{\exp(w_{k0} + \sum_{i=1}^{n} w_{ki} X_i)}{1 + \sum_{j=1}^{R-1} \exp(w_{j0} + \sum_{i=1}^{n} w_{ji} X_i)}$$
  - for k=R $$P(Y = y_R|X) = \frac{1}{1 + \sum_{j=1}^{R-1} \exp(w_{j0} + \sum_{i=1}^{n} w_{ji} X_i)}$$

Training Logistic Regression: MCLE

- We have L training examples: $\{\langle X^1, Y^1\rangle, \ldots \langle X^L, Y^L\rangle\}$

- Maximum likelihood estimate for parameters W

$$W_{MLE} = \arg\max_W P(< X^1, Y^1 > \ldots < X^L, Y^L > | W)$$
$$= \arg\max_W \prod_l P(< X^l, Y^l > | W)$$

- Maximum <u>Conditional</u> Likelihood Estimate (MCLE)

Training Logistic Regression: MCLE

- Choose parameters <w_0 … w_n> to maximize conditional likelihood of training data

$$P(Y=0|X,W) = \frac{1}{1+exp(w_0+\sum_i w_i X_i)}$$

$$P(Y=1|X,W) = \frac{exp(w_0+\sum_i w_i X_i)}{1+exp(w_0+\sum_i w_i X_i)}$$

- Training data D= $\{\langle X^1, Y^1\rangle, \ldots \langle X^L, Y^L\rangle\}$
- Data likelihood = $\prod_l P(X^l, Y^l|W)$
- Data conditional likelihood = $\prod_l P(Y^l|X^l, W)$

$$W_{MCLE} = \arg\max_W \prod_l P(Y^l|W, X^l)$$

Conditional Log Likelihood

$$l(W) \equiv \ln \prod_l P(Y^l|X^l, W) = \sum_l \ln P(Y^l|X^l, W)$$

$$P(Y=0|X,W) = \frac{1}{1+exp(w_0+\sum_i w_i X_i)}$$

$$P(Y=1|X,W) = \frac{exp(w_0+\sum_i w_i X_i)}{1+exp(w_0+\sum_i w_i X_i)}$$

$$\begin{aligned} l(W) &= \sum_l Y^l \ln P(Y^l=1|X^l,W) + (1-Y^l)\ln P(Y^l=0|X^l,W) \\ &= \sum_l Y^l \ln \frac{P(Y^l=1|X^l,W)}{P(Y^l=0|X^l,W)} + \ln P(Y^l=0|X^l,W) \\ &= \sum_l Y^l (w_0 + \sum_i^n w_i X_i^l) - \ln(1+exp(w_0+\sum_i^n w_i X_i^l)) \end{aligned}$$

Maximizing Conditional Log Likelihood

$$P(Y=0|X,W) = \frac{1}{1+exp(w_0+\sum_i w_i X_i)}$$

$$P(Y=1|X,W) = \frac{exp(w_0+\sum_i w_i X_i)}{1+exp(w_0+\sum_i w_i X_i)}$$

$$\begin{aligned} l(W) &\equiv \ln \prod_l P(Y^l|X^l,W) \\ &= \sum_l Y^l(w_0+\sum_i^n w_i X_i^l) - \ln(1+exp(w_0+\sum_i^n w_i X_i^l)) \end{aligned}$$

Good news: l(W) is concave function of W

Bad news: no closed-form solution to maximize l(W)

Gradient Descent

Gradient

$$\nabla E[\vec{w}] \equiv \left[\frac{\partial E}{\partial w_0}, \frac{\partial E}{\partial w_1}, \cdots \frac{\partial E}{\partial w_n}\right]$$

Training rule:

$$\Delta \vec{w} = -\eta \nabla E[\vec{w}]$$

i.e.,

$$\Delta w_i = -\eta \frac{\partial E}{\partial w_i}$$

Maximize Conditional Log Likelihood: Gradient Ascent

$$
\begin{aligned}
l(W) &\equiv \ln \prod_l P(Y^l | X^l, W) \\
&= \sum_l Y^l (w_0 + \sum_i^n w_i X_i^l) - \ln(1 + exp(w_0 + \sum_i^n w_i X_i^l))
\end{aligned}
$$

$$
\frac{\partial l(W)}{\partial w_i} = \sum_l X_i^l (Y^l - \hat{P}(Y^l = 1 | X^l, W))
$$

Maximize Conditional Log Likelihood: Gradient Ascent

$$
\begin{aligned}
l(W) &\equiv \ln \prod_l P(Y^l | X^l, W) \\
&= \sum_l Y^l (w_0 + \sum_i^n w_i X_i^l) - \ln(1 + exp(w_0 + \sum_i^n w_i X_i^l))
\end{aligned}
$$

$$
\frac{\partial l(W)}{\partial w_i} = \sum_l X_i^l (Y^l - \hat{P}(Y^l = 1 | X^l, W))
$$

Gradient ascent algorithm: iterate until change < ε

For all i, repeat

$$
w_i \leftarrow w_i + \eta \sum_l X_i^l (Y^l - \hat{P}(Y^l = 1 | X^l, W))
$$

Logistic Regression: Summary

- Consider learning f: X → Y, where
  - X is a vector of real-valued features, < X_1 … X_n >
  - Y is boolean
  - assume all X_i are conditionally independent given Y
  - model $P(X_i \mid Y = y_k)$ as Gaussian $N(\mu_{ik}, \sigma_i^2)$
  - model P(Y) as Bernoulli (π)
- Then P(Y|X) is of this form and we can directly estimate W

$$P(Y=1|X=<X_1,...X_n>) = \frac{1}{1+exp(w_0+\sum_i w_i X_i)}$$

Linear Discriminant Functions

Augmented Feature Vector

- Linear discriminant function: $\mathbf{g(x) = w^t x + w_0}$
- Can rewrite it:

$$g(x) = \underbrace{\begin{bmatrix} w_0 & w^t \end{bmatrix}}_{\substack{\textit{new weight} \\ \textit{vector a}}} \underbrace{\begin{bmatrix} 1 \\ x \end{bmatrix}}_{\substack{\textit{new feature} \\ \textit{vector y}}} = a^t y = g(y)$$

- $\boldsymbol{y}$ is called the augmented feature vector
- Added a dummy dimension to get a completely equivalent new homogeneous problem

old problem

$$g(x) = w^t x + w_0$$

$$\begin{bmatrix} x_1 \\ \vdots \\ x_d \end{bmatrix}$$

new problem

$$g(y) = a^t y$$

$$\begin{bmatrix} 1 \\ x_1 \\ \vdots \\ x_d \end{bmatrix}$$

- Feature augmentation is done for simpler notation
- From now on, always assume that we have augmented feature vectors
  - Given samples $\mathbf{x}_1,\ldots, \mathbf{x}_n$ convert them to augmented samples $\mathbf{y}_1,\ldots, \mathbf{y}_n$ by adding a new dimension of value 1

$$\mathbf{y}_i = \begin{bmatrix} 1 \\ \mathbf{x}_i \end{bmatrix}$$

Training Error

- For the rest of this part, assume we have 2 classes
  - Samples: $\mathbf{y_1,\ldots, y_n}$, some in class 1, some in class 2
- Use samples to determine weights **a** in the discriminant function $\mathbf{g(y) = a^t y}$
- What should the criterion for determining **a** be?
- For now, suppose we want to minimize the training error (the number of misclassified samples $\mathbf{y_1,\ldots, y_n}$)

- Recall that:

$$g(y_i)>0 \Rightarrow y_i \text{ classified as } c_1$$
$$g(y_i)<0 \Rightarrow y_i \text{ classified as } c_2$$

- Thus training error is **0** if

$$\begin{cases} g(y_i) > 0 & \forall y_i \in c_1 \\ g(y_i) < 0 & \forall y_i \in c_2 \end{cases}$$

"Normalization"

- Thus training error is 0 if: $\begin{cases} a^t y_i > 0 & \forall y_i \in c_1 \\ a^t y_i < 0 & \forall y_i \in c_2 \end{cases}$

- Equivalently, training error is 0 if: $\begin{cases} a^t y_i > 0 & \forall y_i \in c_1 \\ a^t(-y_i) > 0 & \forall y_i \in c_2 \end{cases}$

- This suggests "normalization" (a.k.a. reflection):
  1. Replace all examples from class 2 by:

     $$y_i \to -y_i \qquad \forall y_i \in c_2$$

  2. Seek weight vector $\boldsymbol{a}$ such that

     $$a^t y_i > 0 \qquad \forall y_i$$

  - If such **a** exists, it is called a separating or solution vector
  - Original samples $\mathbf{x}_1, \ldots, \mathbf{x}_n$ can indeed be separated by a line

Normalization

- Seek a hyperplane that separates patterns from different categories
- Seek hyperplane that puts *normalized* patterns on the same (positive) side

Solution Region

- Find weight vector $\boldsymbol{a}$ such that for all samples:

$$\boldsymbol{a}^t \boldsymbol{y}_i = \sum_{k=0}^{d} a_k y_i^{(k)} > 0$$

- In general, there can be many solutions

Solution Region

- Solution region for **a**: set of all possible solutions defined in terms of normal a to the separating hyperplane

Optimization

- Need to minimize a function of many variables

$$J(x) = J(x_1, \ldots, x_d)$$

- We know how to minimize **J(x)**
  - Take partial derivatives and set them to zero

$$\begin{bmatrix} \frac{\partial}{\partial x_1} J(x) \\ \vdots \\ \frac{\partial}{\partial x_d} J(x) \end{bmatrix} = \nabla J(x) = 0$$

gradient

Optimization

- However solving analytically is not always easy
  - For example:

$$\begin{cases} \sin\left(x_1^2 + x_2^3\right) + e^{x_4^2} = 0 \\ \cos\left(x_1^2 + x_2^3\right) + \log\left(x_5^3\right)^{x_4^2} = 0 \end{cases}$$

- Sometimes it is not even possible to write down an analytical expression for the derivative (example later today)

Gradient Descent

- Gradient $\nabla J(x)$ points in direction of steepest increase of **J(x),** and $-\nabla J(x)$ in direction of steepest decrease

Gradient Descent

Gradient Descent for minimizing any function ***J(x)***

- Set k = 1 and $x^{(1)}$ to some initial guess for the weight vector
- While $\eta^{(k)} \left| \nabla J\left(x^{(k)}\right) \right| > \varepsilon$
  - Choose learning rate $\boldsymbol{\eta^{(k)}}$

$$x^{(k+1)} = x^{(k)} - \eta^{(k)} \nabla J(x)$$
$$k = k + 1$$

(update rule)

Gradient Descent

- Gradient decent is guaranteed to only find local minima

- Nevertheless gradient descent is very popular because it is simple and applicable to any function

Gradient Descent

- Main issue: how to set parameter **η** (learning rate)
  - If **η** is too small, too many iterations

  - If **η** is too large may overshoot the minimum and possibly never find it

LDF Criterion Function

- Find weight vector **a** such that for all samples $\mathbf{y}_1,\ldots, \mathbf{y}_n$

$$a^t y_i = \sum_{k=0}^{d} a_k y_i^{(k)} > 0$$

- Need criterion function **J(a)** which is minimized when **a** is a solution vector
- Let $\mathbf{Y_M}$ be the set of examples misclassified by **a**

$$Y_M(a) = \{ y_i \text{ s.t. } a^t y_i < 0 \}$$

- First natural choice: number of misclassified examples

$$J(a) = |Y_M(a)|$$

- Piecewise constant, gradient descent is useless

Perceptron

Perceptron Criterion Function

$$J_p(a) = \sum_{y \in Y_M} \left(-a^t y\right)$$

- If **y** is misclassified, $\mathbf{a^t y < 0}$
- Thus $\mathbf{J_p(a) > 0}$
- $\mathbf{J_p(a)}$ is $\mathbf{||a||}$ times the sum of distances of misclassified examples to decision boundary
- $\mathbf{J_p(a)}$ is piecewise linear and thus suitable for gradient descent

Perceptron Batch Rule

$$J_p(a) = \sum_{y \in Y_M} \left(-a^t y\right)$$

- Gradient of $\mathbf{J_p(a)}$ is: $\nabla J_p(a) = \sum_{y \in Y_M} (-y)$
  - $\mathbf{Y_M}$ are samples misclassified by $\mathbf{a^{(k)}}$
  - It is not possible to solve $\nabla J_p(a) = 0$ analytically because of $\mathbf{Y_M}$
- Update rule for gradient descent: $\mathbf{x}^{(k+1)} = \mathbf{x}^{(k)} - \eta^{(k)} \nabla J(x)$
- Thus the *gradient decent batch update rule for* $\mathbf{J_p(a)}$ is:

$$a^{(k+1)} = a^{(k)} + \eta^{(k)} \sum_{y \in Y_M} y$$

- It is called batch rule because it is based on all misclassified examples

Perceptron Single Sample Rule

- The *gradient decent single sample rule* for $\mathbf{J_p(a)}$ is:

$$\boldsymbol{a}^{(k+1)} = \boldsymbol{a}^{(k)} + \eta^{(k)} \boldsymbol{y}_M$$

  - Note that $\mathbf{y_M}$ is one sample misclassified by **a(k)**
  - Must have a consistent way of visiting samples

- Geometric Interpretation:
  - $\mathbf{y_M}$ misclassified by $\mathbf{a^{(k)}}$ $\left(\boldsymbol{a}^{(k)}\right)^t \boldsymbol{y}_M \leq 0$
  - $\mathbf{y_M}$ is on the wrong side of decision hyperplane
  - Adding $\boldsymbol{\eta y_M}$ to **a** moves the new decision hyperplane in the right direction with respect to $\boldsymbol{y_M}$

Perceptron Single Sample Rule

$$\boldsymbol{a}^{(k+1)} = \boldsymbol{a}^{(k)} + \eta^{(k)} \boldsymbol{y}_M$$

η is too large, previously correctly classified sample $\boldsymbol{y}_k$ is now misclassified

η is too small, $\boldsymbol{y}_M$ is still misclassified

Perceptron Example

	features				grade
name	***good attendance?***	***tall?***	***sleeps in class?***	***chews gum?***	
Jane	*yes (1)*	*yes (1)*	*no (-1)*	*no (-1)*	*A*
Steve	*yes (1)*	*yes (1)*	*yes (1)*	*yes (1)*	*F*
Mary	*no (-1)*	*no (-1)*	*no (-1)*	*yes (1)*	*F*
Peter	*yes (1)*	*no (-1)*	*no (-1)*	*yes (1)*	*A*

- Class 1: students who get A
- Class 2: students who get F

	features					grade
name	***extra***	***good attendance?***	***tall?***	***sleeps in class?***	***chews gum?***	
Jane	*1*	*yes (1)*	*yes (1)*	*no (-1)*	*no (-1)*	*A*
Steve	*1*	*yes (1)*	*yes (1)*	*yes (1)*	*yes (1)*	*F*
Mary	*1*	*no (-1)*	*no (-1)*	*no (-1)*	*yes (1)*	*F*
Peter	*1*	*yes (1)*	*no (-1)*	*no (-1)*	*yes (1)*	*A*

- Augment samples by adding an extra feature (dimension) equal to 1

	features					grade
name	***extra***	***good attendance?***	***tall?***	***sleeps in class?***	***chews gum?***	
Jane	*1*	*yes (1)*	*yes (1)*	*no (-1)*	*no (-1)*	*A*
Steve	*-1*	*yes (-1)*	*yes (-1)*	*yes (-1)*	*yes (-1)*	*F*
Mary	*-1*	*no (1)*	*no (1)*	*no (1)*	*yes (-1)*	*F*
Peter	*1*	*yes (1)*	*no (-1)*	*no (-1)*	*yes (1)*	*A*

- Normalize:
  - Replace all examples from class 2 by their negative values
  - Seek $\boldsymbol{a}$ such that:

$$\boldsymbol{y}_i \rightarrow -\boldsymbol{y}_i \qquad \forall \boldsymbol{y}_i \in \boldsymbol{c}_2$$

$$\boldsymbol{a}^t \boldsymbol{y}_i > 0 \qquad \forall \boldsymbol{y}_i$$

	features					grade
name	***extra***	***good attendance?***	***tall?***	***sleeps in class?***	***chews gum?***	
Jane	*1*	*yes (1)*	*yes (1)*	*no (-1)*	*no (-1)*	*A*
Steve	*-1*	*yes (-1)*	*yes (-1)*	*yes (-1)*	*yes (-1)*	*F*
Mary	*-1*	*no (1)*	*no (1)*	*no (1)*	*yes (-1)*	*F*
Peter	*1*	*yes (1)*	*no (-1)*	*no (-1)*	*yes (1)*	*A*

- Single Sample Rule
  - Sample is misclassified if $\boldsymbol{a}^t \boldsymbol{y}_i = \sum_{k=0}^{4} \boldsymbol{a}_k \boldsymbol{y}_i^{(k)} < 0$
  - Gradient descent single sample rule: $\boldsymbol{a}^{(k+1)} = \boldsymbol{a}^{(k)} + \eta^{(k)} \sum_{\boldsymbol{y} \in Y_M} \boldsymbol{y}$
  - Set **η** fixed learning rate to $\eta^{(k)} = 1$: $\boxed{\boldsymbol{a}^{(k+1)} = \boldsymbol{a}^{(k)} + \boldsymbol{y}_M}$

- Set equal initial weights
  $a^{(1)}$ = [0.25, 0.25, 0.25, 0.25, 0.25]
- Visit all samples sequentially, modifying the weights after each misclassified example

name	$a^t y$	*misclassified?*
Jane	*0.25*1+0.25*1+0.25*1+0.25*(-1)+0.25*(-1) >0*	*no*
Steve	*0.25*(-1)+0.25*(-1)+0.25*(-1)+0.25*(-1)+0.25*(-1)<0*	*yes*

- New weights

$$\begin{aligned} a^{(2)} = a^{(1)} + y_M &= [0.25 \;\; 0.25 \;\; 0.25 \;\; 0.25 \;\; 0.25] + \\ &+ [-1 \;\; -1 \;\; -1 \;\; -1 \;\; -1] = \\ &= [-0.75 \;\; -0.75 \;\; -0.75 \;\; -0.75 \;\; -0.75] \end{aligned}$$

$$a^{(2)} = [-0.75 \;\; -0.75 \;\; -0.75 \;\; -0.75 \;\; -0.75]$$

name	$a^t y$	*misclassified?*
Mary	*-0.75*(-1)-0.75*1 -0.75 *1 -0.75 *1 -0.75*(-1) <0*	*yes*

- New weights

$$a^{(3)} = a^{(2)} + y_M = [-0.75 \;\; -0.75 \;\; -0.75 \;\; -0.75 \;\; -0.75] +$$
$$+ [-1 \;\; 1 \;\; 1 \;\; 1 \;\; -1] =$$
$$= [-1.75 \;\; 0.25 \;\; 0.25 \;\; 0.25 \;\; -1.75]$$

$$a^{(3)} = [-1.75 \quad 0.25 \quad 0.25 \quad 0.25 \quad -1.75]$$

name	$a^t y$	*misclassified?*
Peter	*-1.75 *1 +0.25* 1+0.25* (-1) +0.25 *(-1)-1.75*1 <0*	*yes*

- New weights

$$a^{(4)} = a^{(3)} + y_M = [-1.75 \quad 0.25 \quad 0.25 \quad 0.25 \quad -1.75] +$$
$$+ [1 \quad 1 \quad -1 \quad -1 \quad 1] =$$
$$= [-0.75 \quad 1.25 \quad -0.75 \quad -0.75 \quad -0.75]$$

$$a^{(4)} = [-0.75 \quad 1.25 \quad -0.75 \quad -0.75 \quad -0.75]$$

name	$a^t y$	misclassified?
Jane	-0.75 *1 +1.25*1 -0.75*1 -0.75 *(-1) -0.75 *(-1)+0	no
Steve	-0.75*(-1)+1.25*(-1) -0.75*(-1) -0.75*(-1)-0.75*(-1)>0	no
Mary	-0.75 *(-1)+1.25*1-0.75*1 -0.75 *1 –0.75*(-1) >0	no
Peter	-0.75 *1+ 1.25*1-0.75* (-1)-0.75* (-1) -0.75 *1 >0	no

- Thus the discriminant function is:

$$g(y) = -0.75 * y^{(0)} + 1.25 * y^{(1)} - 0.75 * y^{(2)} - 0.75 * y^{(3)} - 0.75 * y^{(4)}$$

- Converting back to the original features ***x:***

$$g(x) = 1.25 * x^{(1)} - 0.75 * x^{(2)} - 0.75 * x^{(3)} - 0.75 * x^{(4)} - 0.75$$

- Converting back to the original features ***x:***

$$1.25 * x^{(1)} - 0.75 * x^{(2)} - 0.75 * x^{(3)} - 0.75 * x^{(4)} > 0.75 \Rightarrow grade\ A$$
$$1.25 * x^{(1)} - 0.75 * x^{(2)} - 0.75 * x^{(3)} - 0.75 * x^{(4)} < 0.75 \Rightarrow grade\ F$$

good attendance — tall — sleeps in class — chews gum

- This is just one possible solution vector
- If we started with weights

$$\mathbf{a^{(1)}=[0,0.5, 0.5, 0, 0\ ],}$$

- The solution would be [-1,1.5, -0.5, -1, -1]

$$1.5 * x^{(1)} - 0.5 * x^{(2)} - x^{(3)} - x^{(4)} > 1 \Rightarrow grade\ A$$
$$1.5 * x^{(1)} - 0.5 * x^{(2)} - x^{(3)} - x^{(4)} < 1 \Rightarrow grade\ F$$

- In this solution, being tall is the least important feature

LDF: Non-separable Example

- Suppose we have 2 features and the samples are:
  - Class 1: [2,1], [4,3], [3,5]
  - Class 2: [1,3] and [5,6]
- These samples are not separable by a line
- Still would like to get approximate separation by a line
  - A good choice is shown in green
  - Some samples may be “noisy”, and we could accept them being misclassified

LDF: Non-separable Example

- Obtain $\boldsymbol{y_1}, \boldsymbol{y_2}, \boldsymbol{y_3}, \boldsymbol{y_4}$ by adding extra feature and "normalizing"

$$y_1 = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \quad y_2 = \begin{bmatrix} 1 \\ 4 \\ 3 \end{bmatrix} \quad y_3 = \begin{bmatrix} 1 \\ 3 \\ 5 \end{bmatrix} \quad y_4 = \begin{bmatrix} -1 \\ -1 \\ -3 \end{bmatrix} \quad y_5 = \begin{bmatrix} -1 \\ -5 \\ -6 \end{bmatrix}$$

LDF: Non-separable Example

- Apply Perceptron single sample algorithm
- Initial equal weights
  $a^{(1)}$ = [1 1 1]
  - Line equation $x^{(1)}+x^{(2)}+1=0$
- Fixed learning rate ***η = 1***

$$a^{(k+1)} = a^{(k)} + y_M$$

$$y_1 = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \quad y_2 = \begin{bmatrix} 1 \\ 4 \\ 3 \end{bmatrix} \quad y_3 = \begin{bmatrix} 1 \\ 3 \\ 5 \end{bmatrix} \quad y_4 = \begin{bmatrix} -1 \\ -1 \\ -3 \end{bmatrix} \quad y_5 = \begin{bmatrix} -1 \\ -5 \\ -6 \end{bmatrix}$$

- $y^t_1 a^{(1)} = [1\ 1\ 1]*[1\ 2\ 1]^t > 0$ ✓
- $y^t_2 a^{(1)} = [1\ 1\ 1]*[1\ 4\ 3]^t > 0$ ✓
- $y^t_3 a^{(1)} = [1\ 1\ 1]*[1\ 3\ 5]^t > 0$ ✓

LDF: Non-separable Example

$$a^{(1)} = [1\ 1\ 1] \qquad a^{(k+1)} = a^{(k)} + y_M$$

$$y_1 = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \quad y_2 = \begin{bmatrix} 1 \\ 4 \\ 3 \end{bmatrix} \quad y_3 = \begin{bmatrix} 1 \\ 3 \\ 5 \end{bmatrix} \quad y_4 = \begin{bmatrix} -1 \\ -1 \\ -3 \end{bmatrix} \quad y_5 = \begin{bmatrix} -1 \\ -5 \\ -6 \end{bmatrix}$$

- $y^t_4 a^{(1)} = [1\ 1\ 1]*[-1\ -1\ -3]^t = -5 < 0$

$$a^{(2)} = a^{(1)} + y_M = [1\ 1\ 1] + [-1\ -1\ -3] = [0\ 0\ -2]$$

- $y^t_5 a^{(2)} = [0\ 0\ -2]*[-1\ -5\ -6]^t = 12 > 0$ ✔
- $y^t_1 a^{(2)} = [0\ 0\ -2]*[1\ 2\ 1]^t < 0$

$$a^{(3)} = a^{(2)} + y_M = [0\ 0\ -2] + [1\ 2\ 1] = [1\ 2\ -1]$$

LDF: Non-separable Example

$$a^{(3)} = [1\ 2\ -1] \qquad a^{(k+1)} = a^{(k)} + y_M$$

$$y_1 = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \quad y_2 = \begin{bmatrix} 1 \\ 4 \\ 3 \end{bmatrix} \quad y_3 = \begin{bmatrix} 1 \\ 3 \\ 5 \end{bmatrix} \quad y_4 = \begin{bmatrix} -1 \\ -1 \\ -3 \end{bmatrix} \quad y_5 = \begin{bmatrix} -1 \\ -5 \\ -6 \end{bmatrix}$$

- $y^t_2\, a^{(3)} = [1\ 4\ 3]*[1\ 2\ -1]^t = 6 > 0$ ✔
- $y^t_3\, a^{(3)} = [1\ 3\ 5]*[1\ 2\ -1]^t > 0$ ✔
- $y^t_4\, a^{(3)} = [-1\ -1\ -3]*[1\ 2\ -1]^t = 0$

$$a^{(4)} = a^{(3)} + y_M = [1\ 2\ -1] + [-1\ -1\ -3] = [0\ 1\ -4]$$

LDF: Non-separable Example

$$a^{(4)} = [0\ 1\ -4] \qquad a^{(k+1)} = a^{(k)} + y_M$$

$$y_1 = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \quad y_2 = \begin{bmatrix} 1 \\ 4 \\ 3 \end{bmatrix} \quad y_3 = \begin{bmatrix} 1 \\ 3 \\ 5 \end{bmatrix} \quad y_4 = \begin{bmatrix} -1 \\ -1 \\ -3 \end{bmatrix} \quad y_5 = \begin{bmatrix} -1 \\ -5 \\ -6 \end{bmatrix}$$

- $y_5^t a^{(4)}$=[-1 -5 -6]*[0 1 -4]=19>0
- $y_1^t a^{(4)}$=[1 2 1]*[0 1 -4]=-2<0
- ….

LDF: Non-separable Example

- We can continue this forever
- There is no solution vector $\boldsymbol{a}$ satisfying for all $\boldsymbol{i}$

$$\boldsymbol{a}^t \boldsymbol{y}_i = \sum_{k=0}^{5} \boldsymbol{a}_k \boldsymbol{y}_i^{(k)} > 0$$

- Need to stop but at a good point
- Solutions at iterations 900 through 915
  - Some are good and some are not
- How do we stop at a good solution?

Convergence of Perceptron Rules

- If classes are linearly separable and we use fixed learning rate, that is for $\boldsymbol{\eta}^{(k)}$ **=const**
- Then, *both the single sample and batch perceptron rules converge to a correct solution (could be any* ***a*** in the solution space)
- If classes are not linearly separable:
  - The algorithm does not stop, it keeps looking for a solution which does not exist

Convergence of Perceptron Rules

- If classes are not linearly separable:
  - By choosing appropriate learning rate, we can always ensure convergence: $\eta^{(k)} \to 0 \;\; as \;\; k \to \infty$
  - For example inverse linear learning rate: $\eta^{(k)} = \dfrac{\eta^{(1)}}{k}$
  - For inverse linear learning rate, convergence in the linearly separable case can also be proven
  - No guarantee that we stopped at a good point, but there are good reasons to choose inverse linear learning rate